

Assignment 3: Kernel Ridge Regression

CPSC 381/581: Machine Learning

Yale University

Instructor: Alex Wong

Prerequisites:

1. Enable Google Colaboratory as an app on your Google Drive account
2. Create a new Google Colab notebook, this will also create a "Colab Notebooks" directory under "MyDrive" i.e.

```
/content/drive/MyDrive/Colab Notebooks
```

3. Create the following directory structure in your Google Drive

```
/content/drive/MyDrive/Colab Notebooks/CPSC 381-581: Machine Learning/Assignments
```

4. Move the 03_assignment_kernel_regression.ipynb into

```
/content/drive/MyDrive/Colab Notebooks/CPSC 381-581: Machine Learning/Assignments
```

so that its absolute path is

```
/content/drive/MyDrive/Colab Notebooks/CPSC 381-581: Machine Learning/Assignments/03_assignment_kernel_ridge_regression.ipynb
```

In this assignment, we will optimize a kernelized linear function in nonlinear space. We will implement several kernels (linear, polynomial, radial basis function) and train a kernel ridge regression model. We will benchmark our implementation against the one from sci-kit learn, where we should be comparable. Additionally, we will test the speed up for using kernels when the nonlinear mapping function expands feature to high dimensions.

Submission:

1. Implement all TODOs in the code blocks below.
2. Report your training and validation/testing scores.

```
All scores match Sci-kit Learn:
Preprocessing the Diabetes dataset (442 samples, 10 feature
dimensions)
***** Experiments on the Diabetes dataset using linear kernel ridge
regression model with weight decay of 0.01 *****
Results for Scikit-learn model
Training set mean squared error: 25889.5196
Validation set mean squared error: 26708.7982
Testing set mean squared error: 26229.5319
Results for our model
Training set mean squared error: 25889.5196
Validation set mean squared error: 26708.7982
```

```

Testing set mean squared error: 26229.5319
***** Experiments on the Diabetes dataset using polynomial (degree=3)
kernel ridge regression model with weight decay of 0.01 *****
Results for Scikit-learn model
Training set mean squared error: 2458.3355
Validation set mean squared error: 2746.9815
Testing set mean squared error: 3498.1709
Results for our model
Training set mean squared error: 2458.3355
Validation set mean squared error: 2746.9815
Testing set mean squared error: 3498.1709
***** Experiments on the Diabetes dataset using rbf (gamma=1) kernel
ridge regression model with weight decay of 0.01 *****
Results for Scikit-learn model
Training set mean squared error: 2510.6872
Validation set mean squared error: 2697.5951
Testing set mean squared error: 3476.3020
Results for our model
Training set mean squared error: 2510.6872
Validation set mean squared error: 2697.5951
Testing set mean squared error: 3476.3020

```

Preprocessing the Friedman #1 dataset (5000 samples, 20 feature dimensions)

```

***** Experiments on the Friedman #1 dataset using linear kernel ridge
regression model with weight decay of 0.0001 *****
Results for Scikit-learn model
Training set mean squared error: 5.9711
Validation set mean squared error: 6.0310
Testing set mean squared error: 5.9807
Results for our model
Training set mean squared error: 5.9711
Validation set mean squared error: 6.0310
Testing set mean squared error: 5.9807
***** Experiments on the Friedman #1 dataset using polynomial
(degree=3) kernel ridge regression model with weight decay of 0.0001
*****
Results for Scikit-learn model
Training set mean squared error: 0.0354
Validation set mean squared error: 0.2606
Testing set mean squared error: 0.2580
Results for our model
Training set mean squared error: 0.0354
Validation set mean squared error: 0.2606
Testing set mean squared error: 0.2580
***** Experiments on the Friedman #1 dataset using rbf (gamma=1)
kernel ridge regression model with weight decay of 0.0001 *****
Results for Scikit-learn model
Training set mean squared error: 0.0000
Validation set mean squared error: 3.8152
Testing set mean squared error: 3.8152
Results for our model
Training set mean squared error: 0.0000
Validation set mean squared error: 3.8152
Testing set mean squared error: 3.8152

```

AND Kernel ridge regression is faster at 4th and 5th order(see graph at bottom)

Training ridge regression model with degree 2 polynomial expansion

with 5000 samples, 231 feature dimensions

Training time: 4.39ms

Training kernel ridge regression model with degree 2 polynomial with 5000 samples, 20 feature dimensions

Training time: 754.66ms

Training ridge regression model with degree 3 polynomial expansion with 5000 samples, 1771 feature dimensions

Training time: 87.95ms

Training kernel ridge regression model with degree 3 polynomial with 5000 samples, 20 feature dimensions

Training time: 770.12ms

Training ridge regression model with degree 4 polynomial expansion with 5000 samples, 10626 feature dimensions

Training time: 964.51ms

Training kernel ridge regression model with degree 4 polynomial with 5000 samples, 20 feature dimensions

Training time: 768.82ms

Training ridge regression model with degree 5 polynomial expansion with 5000 samples, 53130 feature dimensions

Training time: 3104.30ms

Training kernel ridge regression model with degree 5 polynomial with 5000 samples, 20 feature dimensions

Training time: 799.67ms

3. List any collaborators.

None

IMPORTANT:

- For full credit, your mean squared error for all trained models across all datasets should be same as the scores achieved by sci-kit learn's kernel ridge regression model across training, validation and testing splits. Your kernel ridge regression must be faster than sci-kit linear regression with polynomial expansion at 4th order.

```
In [ ]: import numpy as np
import sklearn.datasets as skdata
import sklearn.metrics as skmetrics
import sklearn.preprocessing as skpreprocess
from sklearn.linear_model import Ridge as RidgeRegressionSciKit
from sklearn.kernel_ridge import KernelRidge as KernelRidgeRegressionSciKit
from matplotlib import pyplot as plt
import warnings
import time

warnings.filterwarnings(action='ignore')
np.random.seed = 1
```

Helper function for plotting

```
In [ ]: def plot_results(axis,
                        x_values,
                        y_values,
                        labels,
                        colors,
                        x_limits,
```

```

        y_limits,
        x_label,
        y_label):
    ...
Plots x and y values using line plot with labels and colors

Args:
    axis : pyplot.ax
            matplotlib subplot axis
    x_values : list[numpy[float32]]
            list of numpy array of x values
    y_values : list[numpy[float32]]
            list of numpy array of y values
    labels : str
            list of names for legend
    colors : str
            colors for each line
    x_limits : list[float32]
            min and max values of x axis
    y_limits : list[float32]
            min and max values of y axis
    x_label : list[float32]
            name of x axis
    y_label : list[float32]
            name of y axis
    ...

# Iterate through x_values, y_values, labels, and colors and plot them
# with associated legend
for x, y, label, color in zip(x_values, y_values, labels, colors):
    axis.plot(x, y, marker='o', color=color, label=label)
    axis.legend(loc='best')

# Set x and y limits
axis.set_xlim(x_limits)
axis.set_ylim(y_limits)

# Set x and y labels
axis.set_xlabel(x_label)
axis.set_ylabel(y_label)

```

Implementation of kernel ridge regression

```

In [ ]: class KernelRidgeRegression(object):

    def __init__(self, kernel_func, degree=None, gamma=None):
        ...
        Class for kernel ridge regression

    Arg(s):
        kernel_func : str
            name of kernel function to use: linear, polynomial, rbf (gaussian)
        degree : int
            p-order for polynomial
        gamma : float
            standard deviation of the Gaussian
        ...

    # Define private variables
    self.__weights = None
    self.__X = None
    self.__kernel_func = kernel_func
    self.__degree = degree
    self.__gamma = gamma

```

```

def __linear_kernel(self, X1, X2):
    """
    Computes the linear kernel function on X1 and X2

    Arg(s):
        X1 : numpy[float32]
            N x d feature vector
        X2 : numpy[float32]
            N x d feature vector
    Returns:
        numpy[float32] : N x N kernel matrix
    """

    # TODO: Implement linear kernel
    return np.dot(X1, X2.T)

def __polynomial_kernel(self, X1, X2, degree):
    """
    Computes the p-order polynomial kernel function on X1 and X2 with c = 1

    Arg(s):
        X1 : numpy[float32]
            N x d feature vector
        X2 : numpy[float32]
            N x d feature vector
        degree : int
            p-order for polynomial
    Returns:
        numpy[float32] : N x N kernel matrix
    """

    c = 1
    return np.power(X1.dot(X2.T) + c, degree)

def __rbf_kernel(self, X1, X2, gamma):
    """
    Computes the RBF (Gaussian) kernel function on X1 and X2

    Arg(s):
        X1 : numpy[float32]
            N x d feature vector
        X2 : numpy[float32]
            N x d feature vector
        gamma : float
            standard deviation of the Gaussian
    Returns:
        numpy[float32] : N x N kernel matrix
    """

    # TODO: Implement RBF kernel using Gaussian form
    diff = X1[:, np.newaxis, :] - X2[np.newaxis, :, :]
    numerator = np.sum(diff ** 2, axis=2)
    return np.exp(-gamma * numerator)

def fit(self, X, y, weight_decay=0):
    """
    Fits the model to X and y via normal equation in kernelized form

    Arg(s):
        X : numpy[float32]
            N x d feature vector
        y : numpy[float32]

```

```

        N x 1 ground-truth label
        weight_decay : float
            weight of weight decay term
    """

    # TODO: Implement the fit function
    self.__X = X.copy()

    if self.__kernel_func == 'linear':
        K = self.__linear_kernel(X, X)

    elif self.__kernel_func == 'polynomial':
        K = self.__polynomial_kernel(X, X, self.__degree)

    elif self.__kernel_func == 'rbf':
        K = self.__rbf_kernel(X, X, self.__gamma)

    else:
        raise ValueError('Unsupported kernel function: {}'.format(self.__kernel_f
A = K + weight_decay * np.eye(K.shape[0])
self.__weights = np.linalg.solve(A, y)

def predict(self, X):
    """
    Predicts the real value for each feature vector X

    Arg(s):
        x : numpy[float32]
            N x d feature vector
    Returns:
        numpy[float32] : N x 1 real value vector ( $\hat{y}$ )
    """

    # TODO: Implement the predict function

    if self.__kernel_func == 'linear':
        K = self.__linear_kernel(X, self.__X)

    elif self.__kernel_func == 'polynomial':
        K = self.__polynomial_kernel(X, self.__X, self.__degree)

    elif self.__kernel_func == 'rbf':
        K = self.__rbf_kernel(X, self.__X, self.__gamma)

    else:
        raise ValueError('Unsupported kernel function: {}'.format(self.__kernel_f

    return np.dot(K, self.__weights)

```

Load datasets

```

In [ ]: # Load diabetes and Friedman #1 datasets
datasets = [
    skdata.load_diabetes(),
    skdata.make_friedman1(n_samples=5000, n_features=20, noise=0.0, random_state=1)
]

dataset_names = [
    'Diabetes',
    'Friedman #1'
]

```

Training, validating and testing kernel ridge regression

```

In [ ]: # Set hyperparameters
diabetes_weight_decay = 1e-2
diabetes_degree = 3
diabetes_gamma = 1

friedman1_weight_decay = 1e-4
friedman1_degree = 3
friedman1_gamma = 1

dataset_hyperparameters = [
    # For diabetes dataset
    [
        diabetes_weight_decay,
        diabetes_degree,
        diabetes_gamma
    ],
    # For Friedman #1 dataset
    [
        friedman1_weight_decay,
        friedman1_degree,
        friedman1_gamma
    ]
]

# Zip up all dataset options
dataset_options = zip(
    datasets,
    dataset_names,
    dataset_hyperparameters)

for options in dataset_options:

    # Unpack dataset options
    dataset, \
        dataset_name, \
        dataset_hyperparameters = options

    weight_decay, degree, gamma = dataset_hyperparameters

    ...
    Create the training, validation and testing splits
    ...
    if dataset_name == 'Friedman #1':
        X, y = dataset
    else:
        X = dataset.data
        y = dataset.target

    print('Preprocessing the {} dataset ({} samples, {} feature dimensions)'.format(d

    # Shuffle the dataset based on sample indices
    shuffled_indices = np.random.permutation(X.shape[0])

    # Choose the first 60% as training set, next 20% as validation and the rest as te
    train_split_idx = int(0.60 * X.shape[0])
    val_split_idx = int(0.80 * X.shape[0])

    train_indices = shuffled_indices[0:train_split_idx]
    val_indices = shuffled_indices[train_split_idx:val_split_idx]
    test_indices = shuffled_indices[val_split_idx:]

    # Select the examples from X and y to construct our training, validation, testing
    X_train, y_train = X[train_indices, :], y[train_indices]

```

```

X_val, y_val = X[val_indices, :], y[val_indices]
X_test, y_test = X[test_indices, :], y[test_indices]

for kernel in ['linear', 'polynomial', 'rbf']:

    """
    Trains and tests kernel ridge regression model for different kernels
    """

    if kernel == 'linear':
        print('***** Experiments on the {} dataset using {} kernel ridge regression model with {}
              dataset_name,
              kernel,
              weight_decay))

        # TODO: Instantiate KernelRidgeRegressionSciKit with linear kernel
        model_scikit = KernelRidgeRegressionSciKit(kernel=kernel, alpha=weight_decay)

        # TODO: Instantiate our kernel ridge regression model with linear kernel
        model_ours = KernelRidgeRegression('linear')

    elif kernel == 'polynomial':
        print('***** Experiments on the {} dataset using {} (degree={}) kernel ridge regression model with {}
              dataset_name,
              kernel,
              degree,
              weight_decay))

        # TODO: Instantiate KernelRidgeRegressionSciKit with a polynomial kernel
        model_scikit = KernelRidgeRegressionSciKit(kernel=kernel, alpha=weight_decay)

        # TODO: Instantiate our kernel ridge regression model with polynomial kernel
        model_ours = KernelRidgeRegression('polynomial', degree=degree)

    elif kernel == 'rbf':
        print('***** Experiments on the {} dataset using {} (gamma={}) kernel ridge regression model with {}
              dataset_name,
              kernel,
              gamma,
              weight_decay))

        # TODO: Instantiate KernelRidgeRegressionSciKit with an rbf kernel. Please use gamma parameter
        model_scikit = KernelRidgeRegressionSciKit(kernel=kernel, alpha=weight_decay)

        # TODO: Instantiate our kernel ridge regression model with an rbf kernel
        model_ours = KernelRidgeRegression('rbf', gamma=gamma)
    else:
        raise ValueError('Unsupported kernel function: {}'.format(kernel))

print('Results for Scikit-learn model')

# TODO: Train scikit-learn model
model_scikit.fit(X_train, y_train)

# TODO: Score model using mean squared error on training set
predictions_scikit_train = model_scikit.predict(X_train)
mse_scikit_train = skmetrics.mean_squared_error(y_train, predictions_scikit_train)
print('Training set mean squared error: {:.4f}'.format(mse_scikit_train))

# TODO: Score model using mean squared error validation set

mse_scikit_val = skmetrics.mean_squared_error(y_val, model_scikit.predict(X_val))
print('Validation set mean squared error: {:.4f}'.format(mse_scikit_val))

# TODO: Score model using mean squared error testing set

```



```
mse_scikit_test = skmetrics.mean_squared_error(y_test, model_scikit.predict(X_test))
print('Testing set mean squared error: {:.4f}'.format(mse_scikit_test))

print('Results for our model')

# TODO: Train our model
model_ours.fit(X_train, y_train, weight_decay=0.01)

# TODO: Score model using mean squared error on training set

mse_ours_train = skmetrics.mean_squared_error(np.squeeze(y_train), np.squeeze(model_ours.predict(X_train)))
print('Training set mean squared error: {:.4f}'.format(mse_ours_train))

# TODO: Score model using mean squared error validation set

mse_ours_val = skmetrics.mean_squared_error(np.squeeze(y_val), np.squeeze(model_ours.predict(X_val)))
print('Validation set mean squared error: {:.4f}'.format(mse_ours_val))

# TODO: Score model using mean squared error testing set

mse_ours_test = skmetrics.mean_squared_error(np.squeeze(y_test), np.squeeze(model_ours.predict(X_test)))
print('Testing set mean squared error: {:.4f}'.format(mse_ours_test))

print('')
```

```

Preprocessing the Diabetes dataset (442 samples, 10 feature dimensions)
***** Experiments on the Diabetes dataset using linear kernel ridge regression model with weight decay of 0.01 *****
Results for Scikit-learn model
Training set mean squared error: 26221.6241
Validation set mean squared error: 26310.9583
Testing set mean squared error: 26436.4242
Results for our model
Training set mean squared error: 26221.6241
Validation set mean squared error: 26310.9583
Testing set mean squared error: 26436.4242
***** Experiments on the Diabetes dataset using polynomial (degree=3) kernel ridge regression model with weight decay of 0.01 *****
Results for Scikit-learn model
Training set mean squared error: 2402.3747
Validation set mean squared error: 3476.8757
Testing set mean squared error: 3023.2954
Results for our model
Training set mean squared error: 2402.3747
Validation set mean squared error: 3476.8757
Testing set mean squared error: 3023.2954
***** Experiments on the Diabetes dataset using rbf (gamma=1) kernel ridge regression model with weight decay of 0.01 *****
Results for Scikit-learn model
Training set mean squared error: 2442.4863
Validation set mean squared error: 3438.5361
Testing set mean squared error: 3020.9325
Results for our model
Training set mean squared error: 2442.4863
Validation set mean squared error: 3438.5361
Testing set mean squared error: 3020.9325

Preprocessing the Friedman #1 dataset (5000 samples, 20 feature dimensions)
***** Experiments on the Friedman #1 dataset using linear kernel ridge regression model with weight decay of 0.0001 *****
Results for Scikit-learn model
Training set mean squared error: 6.3020
Validation set mean squared error: 5.5840
Testing set mean squared error: 5.4993
Results for our model
Training set mean squared error: 6.3020
Validation set mean squared error: 5.5840
Testing set mean squared error: 5.4993
***** Experiments on the Friedman #1 dataset using polynomial (degree=3) kernel ridge regression model with weight decay of 0.0001 *****
Results for Scikit-learn model
Training set mean squared error: 0.0383
Validation set mean squared error: 0.2762
Testing set mean squared error: 0.2630
Results for our model
Training set mean squared error: 0.0383
Validation set mean squared error: 0.2762
Testing set mean squared error: 0.2630
***** Experiments on the Friedman #1 dataset using rbf (gamma=1) kernel ridge regression model with weight decay of 0.0001 *****
Results for Scikit-learn model
Training set mean squared error: 0.0000
Validation set mean squared error: 4.0765
Testing set mean squared error: 3.6517
Results for our model
Training set mean squared error: 0.0000
Validation set mean squared error: 4.0765
Testing set mean squared error: 3.6517

```

Comparing run time for polynomial kernel and polynomial feature expansion

```

In [ ]: # Define weight decay and polynomial degrees
weight_decay = 1

degrees = [
    2, 3, 4, 5
]

# Lists to hold time elapsed for
times_elapsed_poly_expand = []
times_elapsed_poly_kernel = []

# Select Friedman #1 dataset
dataset = skdata.make_friedman1(n_samples=5000, n_features=20, noise=1.0, random_stat

X, y = dataset

for degree in degrees:

    # TODO: Initialize polynomial expansion
    poly_transform = skpreprocess.PolynomialFeatures(degree=degree)

    # TODO: Compute the polynomial terms needed for the data
    poly_transform.fit(X)

    # TODO: Transform the data by nonlinear mapping
    X_poly = poly_transform.transform(X)

    # TODO: Initialize sci-kit ridge regression model
    model_poly_expand = RidgeRegressionSciKit(alpha=weight_decay)

    print('Training ridge regression model with degree {} polynomial expansion with {
        degree,
        X_poly.shape[0],
        X_poly.shape[1])

    time_start = time.time()

    # TODO: Train sci-kit ridge regression model on polynomial expanded X
    model_poly_expand.fit(X_poly, y)

    time_elapsed_poly_expand = 1000 * (time.time() - time_start)

    print('Training time: {:.2f}ms'.format(time_elapsed_poly_expand))

    # TODO: Append training time to list of time elapsed for polynomial feature expan
    times_elapsed_poly_expand.append(time_elapsed_poly_expand)

    # TODO: Initialize our polynomial kernel ridge regression model
    model_poly_kernel = KernelRidgeRegression('polynomial', degree=degree)

    print('Training kernel ridge regression model with degree {} polynomial with {} s
        degree,
        X.shape[0],
        X.shape[1])

    time_start = time.time()

    # TODO: Train our polynomial kernel ridge regression model on X
    model_poly_kernel.fit(X, y, weight_decay)

    time_elapsed_poly_kernel = 1000 * (time.time() - time_start)

    print('Training time: {:.2f}ms'.format(time_elapsed_poly_kernel))

```

```

# TODO: Append training time to list of time elapsed for polynomial kernel
times_elapsed_poly_kernel.append(time_elapsed_poly_kernel)

print('')

# Create figure for training, validation and testing scores for different features
labels = ['Polynomial Expansion', 'Polynomial Kernel']
colors = ['blue', 'red']

# TODO: Create a subplot of a 1 by 1 figure to plot MSE for training and testing
fig = plt.figure(figsize=(10, 5))
ax = fig.add_subplot(1, 1, 1)

# TODO: Set x values (polynomial degree) and y values (time in ms in log scale)
x_values = [degrees] * 2
y_values = [times_elapsed_poly_expand, times_elapsed_poly_kernel]
y_values = [np.log(y) for y in y_values]

# TODO: Plot MSE scores for training and testing sets
# Set x limits between 1 to 1 + maximum value of all degrees and y limits between 1 and
# Set x label to degree and y label to milliseconds (log scale)' and y label to 'MSE'
plot_results(
    axis=ax,
    x_values=x_values,
    y_values=y_values,
    labels=labels,
    colors=colors,
    x_limits=[1, max(degrees) + 1],
    y_limits=[1, max(max(y_values[0]), max(y_values[1])) + 1],
    x_label='Degree',
    y_label='Milliseconds (log scale)')

# TODO: Create plot title of 'Comparing Polynomial Feature Expansion and Polynomial Kernel Run-times'
plt.title('Comparing Polynomial Feature Expansion and Polynomial Kernel Run-times')

```

Training ridge regression model with degree 2 polynomial expansion with 5000 samples, 231 feature dimensions
Training time: 5.46ms
Training kernel ridge regression model with degree 2 polynomial with 5000 samples, 20 feature dimensions
Training time: 729.92ms

Training ridge regression model with degree 3 polynomial expansion with 5000 samples, 1771 feature dimensions
Training time: 87.50ms
Training kernel ridge regression model with degree 3 polynomial with 5000 samples, 20 feature dimensions
Training time: 747.73ms

Training ridge regression model with degree 4 polynomial expansion with 5000 samples, 10626 feature dimensions
Training time: 995.10ms
Training kernel ridge regression model with degree 4 polynomial with 5000 samples, 20 feature dimensions
Training time: 812.18ms

Training ridge regression model with degree 5 polynomial expansion with 5000 samples, 53130 feature dimensions
Training time: 3843.72ms
Training kernel ridge regression model with degree 5 polynomial with 5000 samples, 20 feature dimensions
Training time: 797.31ms

Out[]: Text(0.5, 1.0, 'Comparing Polynomial Feature Expansion and Polynomial Kernel Run-times')

