

Exercise 4: Gradient Descent for Linear Regression

CPSC 381/581: Machine Learning

Yale University

Instructor: Alex Wong

Prerequisites:

1. Enable Google Colaboratory as an app on your Google Drive account
2. Create a new Google Colab notebook, this will also create a "Colab Notebooks" directory under "MyDrive" i.e.

```
/content/drive/MyDrive/Colab Notebooks
```

3. Create the following directory structure in your Google Drive

```
/content/drive/MyDrive/Colab Notebooks/CPSC 381-581: Machine  
Learning/Exercises
```

4. Move the 04_exercise_gradient_descent.ipynb into

```
/content/drive/MyDrive/Colab Notebooks/CPSC 381-581: Machine  
Learning/Exercises
```

so that its absolute path is

```
/content/drive/MyDrive/Colab Notebooks/CPSC 381-581: Machine  
Learning/Exercises/04_exercise_gradient_descent.ipynb
```

In this exercise, we will optimize a linear function for the regression task using the gradient descent for mean squared and half mean squared losses. We will test them on several datasets.

Submission:

1. Implement all TODOs in the code blocks below.
2. Report your training, validation, and testing scores.

Report validation and testing scores here.

For full credit, your mean squared error scores for models optimized using `mean_squared` and `half_mean_squared` losses on Diabetes dataset should be no more than 15% worse the mean squared error scores achieved by `sci-kit learn`'s linear regression model across training, validation and testing splits. Your mean squared error scores on California housing price dataset should be no more than 20% worse.

3. List any collaborators.

Collaborators: Doe, Jane (Please write names in <Last Name, First Name> format)

Collaboration details: Discussed ... implementation details with Jane Doe.

Import packages

```
In [ ]: import numpy as np
import sklearn.datasets as skdata
import sklearn.metrics as skmetrics
from sklearn.linear_model import LinearRegression as LinearRegressionSciKit
import warnings

warnings.filterwarnings(action='ignore')
np.random.seed = 1
```

Implementation of our Gradient Descent optimizer for mean squared and half mean squared loss

```
In [ ]: class GradientDescentOptimizer(object):

    def __init__(self):
        pass

    def _compute_gradients(self, w, x, y, loss_func):
        """
        Returns the gradient of mean squared or half mean squared loss

        Arg(s):
            w : numpy[float32]
                d x 1 weight vector
            x : numpy[float32]
                d x N feature vector
            y : numpy[float32]
                1 x N groundtruth vector
            loss_func : str
                loss type either 'mean_squared', or 'half_mean_squared'
        Returns:
            numpy[float32] : d x 1 gradients
        """

        # TODO: Implements the _compute_gradients function
        if loss_func == 'mean_squared':
            gradients = (np.matmul(w.T, x) - y) * x

            return 2.0 * np.mean(gradients, axis=1, keepdims=True)
        elif loss_func == 'half_mean_squared':
            gradients = (np.matmul(w.T, x) - y) * x

            return np.mean(gradients, axis=1, keepdims=True)
        else:
            raise ValueError('Unsupported loss function: {}'.format(loss_func))

    def update(self, w, x, y, alpha, loss_func):
        """
        Updates the weight vector based on mean squared or half mean squared loss

        Arg(s):
            w : numpy[float32]
                d x 1 weight vector
            x : numpy[float32]
                d x N feature vector
            y : numpy[float32]
                1 x N groundtruth vector
            alpha : float
                learning rate
```

```

        loss_func : str
            loss type either 'mean_squared', or 'half_mean_squared'
Returns:
    numpy[float32] : d x 1 weights
    ...

# TODO: Implement the optimizer update function

    return w - alpha * self._compute_gradients(w, x, y, loss_func)

```

Implementation of Linear Regression with Gradient Descent optimizer

```

In [ ]: class LinearRegressionGradientDescent(object):

    def __init__(self):
        # Define private variables
        self.__weights = None
        self.__optimizer = GradientDescentOptimizer()

    def fit(self, x, y, T, alpha, loss_func='mean_squared'):
        ...

        Fits the model to x and y by updating the weight vector
        using gradient descent

        Arg(s):
            x : numpy[float32]
                d x N feature vector
            y : numpy[float32]
                1 x N groundtruth vector
            T : int
                number of iterations to train
            alpha : float
                learning rate
            loss_func : str
                loss function to use
        ...

        # TODO: Implement the fit function
        self.__weights = np.zeros([x.shape[0], 1])

        for t in range(1, T + 1):

            # TODO: Compute loss function
            loss = self._compute_loss(
                x=x,
                y=y,
                loss_func=loss_func)

            if (t % 10000) == 0:
                print('Step={} Loss={:.4f}'.format(t, loss))

            # TODO: Update weights
            self.__weights = self.__optimizer.update(
                w=self.__weights,
                x=x,
                y=y,
                alpha=alpha,
                loss_func=loss_func)

    def predict(self, x):
        ...

        Predicts the label for each feature vector x

        Arg(s):

```

```

        x : numpy[float32]
            d x N feature vector
Returns:
    numpy[float32] : 1 x N vector
    ...

# TODO: Implements the predict function

    return np.matmul(self.__weights.T, x)

def _compute_loss(self, x, y, loss_func):
    """
    Returns the gradient of the mean squared or half mean squared loss

    Arg(s):
        x : numpy[float32]
            d x N feature vector
        y : numpy[float32]
            1 x N groundtruth vector
        loss_func : str
            loss type either 'mean_squared', or 'half_mean_squared'
    Returns:
        float : loss
    """

    # TODO: Implements the _compute_loss function
    predictions = self.predict(x)

    if loss_func == 'mean_squared':
        # TODO: Implements loss for mean squared loss
        loss = np.mean((predictions - y) ** 2)
    elif loss_func == 'half_mean_squared':
        # TODO: Implements loss for half mean squared loss
        loss = 0.50 * np.mean((predictions - y) ** 2)
    else:
        raise ValueError('Unsupported loss function: {}'.format(loss_func))

    return loss

```

Implementing training and validation loop for linear regression

```

In [ ]: # Load Diabetes and California housing prices dataset
datasets = [
    skdata.load_diabetes(),
    skdata.fetch_california_housing()
]
dataset_names = [
    'Diabetes',
    'California housing prices'
]

# Loss functions to minimize
dataset_loss_funcs = [
    ['mean_squared', 'half_mean_squared'],
    ['mean_squared', 'half_mean_squared']
]

# TODO: Select learning rates (alpha) for mean squared and half mean squared loss
dataset_alphas = [
    [1, 1],
    [1e-7, 2.5e-7]
]

# TODO: Select number of steps (T) to train for mean squared and half mean squared lo
dataset_Ts = [

```

```

[100000, 100000],
[2000000, 100000]
]

for dataset_options in zip(datasets, dataset_names, dataset_loss_funcs, dataset_alpha

    dataset, dataset_name, loss_funcs, alphas, Ts = dataset_options

    ...
    Create the training, validation and testing splits
    ...
    x = dataset.data
    y = dataset.target

    # Shuffle the dataset based on sample indices
    shuffled_indices = np.random.permutation(x.shape[0])

    # Choose the first 80% as training set, next 10% as validation and the rest as te
    train_split_idx = int(0.80 * x.shape[0])
    val_split_idx = int(0.90 * x.shape[0])

    train_indices = shuffled_indices[0:train_split_idx]
    val_indices = shuffled_indices[train_split_idx:val_split_idx]
    test_indices = shuffled_indices[val_split_idx:]

    # Select the examples from x and y to construct our training, validation, testing
    x_train, y_train = x[train_indices, :], y[train_indices]
    x_val, y_val = x[val_indices, :], y[val_indices]
    x_test, y_test = x[test_indices, :], y[test_indices]

    ...
    Trains and tests Linear Regression model from scikit-learn
    ...
    # TODO: Initialize scikit-learn linear regression model without bias
    model_scikit = LinearRegressionSciKit(fit_intercept=False)

    # TODO: Trains scikit-learn linear regression model

    model_scikit.fit(x_train, y_train)

    print('***** Results of scikit-learn linear regression model on {} dataset *****'
          dataset_name))

    # TODO: Test model on training set
    predictions_train = model_scikit.predict(x_train)

    score_mse_train = skmetrics.mean_squared_error(y_train, predictions_train)
    print('Training set mean squared error: {:.4f}'.format(score_mse_train))

    score_r2_train = skmetrics.r2_score(y_train, predictions_train)
    print('Training set r-squared scores: {:.4f}'.format(score_r2_train))

    # TODO: Test model on validation set
    predictions_val = model_scikit.predict(x_val)

    score_mse_val = skmetrics.mean_squared_error(y_val, predictions_val)
    print('Validation set mean squared error: {:.4f}'.format(score_mse_val))

    score_r2_val = skmetrics.r2_score(y_val, predictions_val)
    print('Validation set r-squared scores: {:.4f}'.format(score_r2_val))

    # TODO: Test model on testing set
    predictions_test = model_scikit.predict(x_test)

    score_mse_test = skmetrics.mean_squared_error(y_test, predictions_test)

```

```

print('Testing set mean squared error: {:.4f}'.format(score_mse_test))

score_r2_test = skmetrics.r2_score(y_test, predictions_test)
print('Testing set r-squared scores: {:.4f}'.format(score_r2_test))

'''
Trains and tests our linear regression model using different solvers
'''

# Take the transpose of the dataset to match the dimensions discussed in lecture
# i.e., (N x d) to (d x N)
x_train = np.transpose(x_train, axes=(1, 0))
x_val = np.transpose(x_val, axes=(1, 0))
x_test = np.transpose(x_test, axes=(1, 0))
y_train = np.expand_dims(y_train, axis=0)
y_val = np.expand_dims(y_val, axis=0)
y_test = np.expand_dims(y_test, axis=0)

for loss_func, alpha, T in zip(loss_funcs, alphas, Ts):

    # TODO: Initialize our linear regression model
    model_ours = LinearRegressionGradientDescent()

    print('***** Results of our linear regression model trained with {} loss, alp
          loss_func, alpha, T, dataset_name))

    # TODO: Train model on training set
    model_ours.fit(x_train, y_train, T, alpha, loss_func)
    # TODO: Make predictions
    predictions_train = model_ours.predict(x_train)

    # TODO: Test model on training set using mean squared error and r-squared
    score_mse_train = model_ours._compute_loss(x_train, y_train, loss_func)
    print('Training set mean squared error: {:.4f}'.format(score_mse_train))
    score_r2_train = skmetrics.r2_score(np.squeeze(y_train), np.squeeze(predictions_train))
    print('Training set r-squared scores: {:.4f}'.format(score_r2_train))

    # TODO: Test model on validation set using mean squared error and r-squared
    predictions_val = model_ours.predict(x_val)

    score_mse_val = model_ours._compute_loss(x_val, y_val, loss_func)
    print('Validation set mean squared error: {:.4f}'.format(score_mse_val))
    score_r2_val = skmetrics.r2_score(np.squeeze(y_val), np.squeeze(predictions_val))

    print('Validation set r-squared scores: {:.4f}'.format(score_r2_val))

    # TODO: Test model on testing set using mean squared error and r-squared
    predictions_test = model_ours.predict(x_test)

    score_mse_test = model_ours._compute_loss(x_test, y_test, loss_func)
    print('Testing set mean squared error: {:.4f}'.format(score_mse_test))
    score_r2_test = skmetrics.r2_score(np.squeeze(y_test), np.squeeze(predictions_test))
    print('Testing set r-squared scores: {:.4f}'.format(score_r2_test))

```

```

***** Results of scikit-learn linear regression model on Diabetes dataset *****
Training set mean squared error: 25924.5928
Training set r-squared scores: -3.2354
Validation set mean squared error: 28170.6028
Validation set r-squared scores: -4.0233
Testing set mean squared error: 27257.1553
Testing set r-squared scores: -4.9347
***** Results of our linear regression model trained with mean_squared loss, alpha=1 and T=100000 on Diabetes dataset *****
Step=10000 Loss=25940.3928
Step=20000 Loss=25931.6810
Step=30000 Loss=25927.7733
Step=40000 Loss=25926.0199
Step=50000 Loss=25925.2332
Step=60000 Loss=25924.8801
Step=70000 Loss=25924.7217
Step=80000 Loss=25924.6507
Step=90000 Loss=25924.6188
Step=100000 Loss=25924.6045
Training set mean squared error: 25924.6045
Training set r-squared scores: -3.2354
Validation set mean squared error: 28162.2910
Validation set r-squared scores: -4.0219
Testing set mean squared error: 27261.0160
Testing set r-squared scores: -4.9355
***** Results of our linear regression model trained with half_mean_squared loss, alpha=1 and T=100000 on Diabetes dataset *****
Step=10000 Loss=12974.1281
Step=20000 Loss=12970.1962
Step=30000 Loss=12967.5872
Step=40000 Loss=12965.8404
Step=50000 Loss=12964.6704
Step=60000 Loss=12963.8866
Step=70000 Loss=12963.3616
Step=80000 Loss=12963.0100
Step=90000 Loss=12962.7744
Step=100000 Loss=12962.6166
Training set mean squared error: 12962.6166
Training set r-squared scores: -3.2355
Validation set mean squared error: 14054.8160
Validation set r-squared scores: -4.0125
Testing set mean squared error: 13643.0420
Testing set r-squared scores: -4.9410
***** Results of scikit-learn linear regression model on California housing prices dataset *****
Training set mean squared error: 0.5998
Training set r-squared scores: 0.5481
Validation set mean squared error: 0.6166
Validation set r-squared scores: 0.5598
Testing set mean squared error: 0.6364
Testing set r-squared scores: 0.5082
***** Results of our linear regression model trained with mean_squared loss, alpha=1e-07 and T=2000000 on California housing prices dataset *****
Step=10000 Loss=1.3036
Step=20000 Loss=1.2917
Step=30000 Loss=1.2805
Step=40000 Loss=1.2698
Step=50000 Loss=1.2594
Step=60000 Loss=1.2494
Step=70000 Loss=1.2396
Step=80000 Loss=1.2299
Step=90000 Loss=1.2205
Step=100000 Loss=1.2113
Step=110000 Loss=1.2022
Step=120000 Loss=1.1933
Step=130000 Loss=1.1846

```

Step=140000	Loss=1.1760
Step=150000	Loss=1.1675
Step=160000	Loss=1.1592
Step=170000	Loss=1.1510
Step=180000	Loss=1.1430
Step=190000	Loss=1.1351
Step=200000	Loss=1.1274
Step=210000	Loss=1.1198
Step=220000	Loss=1.1123
Step=230000	Loss=1.1049
Step=240000	Loss=1.0977
Step=250000	Loss=1.0906
Step=260000	Loss=1.0836
Step=270000	Loss=1.0767
Step=280000	Loss=1.0700
Step=290000	Loss=1.0633
Step=300000	Loss=1.0568
Step=310000	Loss=1.0504
Step=320000	Loss=1.0441
Step=330000	Loss=1.0379
Step=340000	Loss=1.0318
Step=350000	Loss=1.0258
Step=360000	Loss=1.0199
Step=370000	Loss=1.0141
Step=380000	Loss=1.0084
Step=390000	Loss=1.0027
Step=400000	Loss=0.9972
Step=410000	Loss=0.9918
Step=420000	Loss=0.9864
Step=430000	Loss=0.9812
Step=440000	Loss=0.9760
Step=450000	Loss=0.9709
Step=460000	Loss=0.9659
Step=470000	Loss=0.9610
Step=480000	Loss=0.9562
Step=490000	Loss=0.9514
Step=500000	Loss=0.9467
Step=510000	Loss=0.9421
Step=520000	Loss=0.9376
Step=530000	Loss=0.9331
Step=540000	Loss=0.9287
Step=550000	Loss=0.9244
Step=560000	Loss=0.9201
Step=570000	Loss=0.9159
Step=580000	Loss=0.9118
Step=590000	Loss=0.9078
Step=600000	Loss=0.9038
Step=610000	Loss=0.8998
Step=620000	Loss=0.8960
Step=630000	Loss=0.8922
Step=640000	Loss=0.8884
Step=650000	Loss=0.8847
Step=660000	Loss=0.8811
Step=670000	Loss=0.8775
Step=680000	Loss=0.8740
Step=690000	Loss=0.8706
Step=700000	Loss=0.8671
Step=710000	Loss=0.8638
Step=720000	Loss=0.8605
Step=730000	Loss=0.8572
Step=740000	Loss=0.8540
Step=750000	Loss=0.8509
Step=760000	Loss=0.8478
Step=770000	Loss=0.8447
Step=780000	Loss=0.8417
Step=790000	Loss=0.8387

Step=800000	Loss=0.8358
Step=810000	Loss=0.8329
Step=820000	Loss=0.8301
Step=830000	Loss=0.8273
Step=840000	Loss=0.8246
Step=850000	Loss=0.8219
Step=860000	Loss=0.8192
Step=870000	Loss=0.8166
Step=880000	Loss=0.8140
Step=890000	Loss=0.8115
Step=900000	Loss=0.8089
Step=910000	Loss=0.8065
Step=920000	Loss=0.8040
Step=930000	Loss=0.8016
Step=940000	Loss=0.7993
Step=950000	Loss=0.7970
Step=960000	Loss=0.7947
Step=970000	Loss=0.7924
Step=980000	Loss=0.7902
Step=990000	Loss=0.7880
Step=1000000	Loss=0.7858
Step=1010000	Loss=0.7837
Step=1020000	Loss=0.7816
Step=1030000	Loss=0.7795
Step=1040000	Loss=0.7775
Step=1050000	Loss=0.7755
Step=1060000	Loss=0.7735
Step=1070000	Loss=0.7716
Step=1080000	Loss=0.7696
Step=1090000	Loss=0.7677
Step=1100000	Loss=0.7659
Step=1110000	Loss=0.7640
Step=1120000	Loss=0.7622
Step=1130000	Loss=0.7604
Step=1140000	Loss=0.7587
Step=1150000	Loss=0.7569
Step=1160000	Loss=0.7552
Step=1170000	Loss=0.7535
Step=1180000	Loss=0.7519
Step=1190000	Loss=0.7502
Step=1200000	Loss=0.7486
Step=1210000	Loss=0.7470
Step=1220000	Loss=0.7455
Step=1230000	Loss=0.7439
Step=1240000	Loss=0.7424
Step=1250000	Loss=0.7409
Step=1260000	Loss=0.7394
Step=1270000	Loss=0.7379
Step=1280000	Loss=0.7365
Step=1290000	Loss=0.7351
Step=1300000	Loss=0.7337
Step=1310000	Loss=0.7323
Step=1320000	Loss=0.7309
Step=1330000	Loss=0.7296
Step=1340000	Loss=0.7283
Step=1350000	Loss=0.7270
Step=1360000	Loss=0.7257
Step=1370000	Loss=0.7244
Step=1380000	Loss=0.7231
Step=1390000	Loss=0.7219
Step=1400000	Loss=0.7207
Step=1410000	Loss=0.7195
Step=1420000	Loss=0.7183
Step=1430000	Loss=0.7171
Step=1440000	Loss=0.7160
Step=1450000	Loss=0.7148

```
Step=1460000 Loss=0.7137
Step=1470000 Loss=0.7126
Step=1480000 Loss=0.7115
Step=1490000 Loss=0.7104
Step=1500000 Loss=0.7094
Step=1510000 Loss=0.7083
Step=1520000 Loss=0.7073
Step=1530000 Loss=0.7063
Step=1540000 Loss=0.7053
Step=1550000 Loss=0.7043
Step=1560000 Loss=0.7033
Step=1570000 Loss=0.7023
Step=1580000 Loss=0.7014
Step=1590000 Loss=0.7005
Step=1600000 Loss=0.6995
Step=1610000 Loss=0.6986
Step=1620000 Loss=0.6977
Step=1630000 Loss=0.6968
Step=1640000 Loss=0.6959
Step=1650000 Loss=0.6951
Step=1660000 Loss=0.6942
Step=1670000 Loss=0.6934
Step=1680000 Loss=0.6925
Step=1690000 Loss=0.6917
Step=1700000 Loss=0.6909
Step=1710000 Loss=0.6901
Step=1720000 Loss=0.6893
Step=1730000 Loss=0.6885
Step=1740000 Loss=0.6878
Step=1750000 Loss=0.6870
Step=1760000 Loss=0.6863
Step=1770000 Loss=0.6855
Step=1780000 Loss=0.6848
Step=1790000 Loss=0.6841
Step=1800000 Loss=0.6834
Step=1810000 Loss=0.6827
Step=1820000 Loss=0.6820
Step=1830000 Loss=0.6813
Step=1840000 Loss=0.6806
Step=1850000 Loss=0.6799
Step=1860000 Loss=0.6793
Step=1870000 Loss=0.6786
Step=1880000 Loss=0.6780
Step=1890000 Loss=0.6774
Step=1900000 Loss=0.6768
Step=1910000 Loss=0.6761
Step=1920000 Loss=0.6755
Step=1930000 Loss=0.6749
Step=1940000 Loss=0.6743
Step=1950000 Loss=0.6738
Step=1960000 Loss=0.6732
Step=1970000 Loss=0.6726
Step=1980000 Loss=0.6720
Step=1990000 Loss=0.6715
Step=2000000 Loss=0.6709
```

Training set mean squared error: 0.6709

Training set r-squared scores: 0.4946

Validation set mean squared error: 0.6986

Validation set r-squared scores: 0.5013

Testing set mean squared error: 0.6648

Testing set r-squared scores: 0.4863

***** Results of our linear regression model trained with half_mean_squared loss, alpha=2.5e-07 and T=100000 on California housing prices dataset *****

```
Step=10000 Loss=0.6503
```

```
Step=20000 Loss=0.6430
```

```
Step=30000 Loss=0.6362
```

```
Step=40000 Loss=0.6297
Step=50000 Loss=0.6234
Step=60000 Loss=0.6174
Step=70000 Loss=0.6114
Step=80000 Loss=0.6056
Step=90000 Loss=0.6000
Step=100000 Loss=0.5945
Training set mean squared error: 0.5945
Training set r-squared scores: 0.1044
Validation set mean squared error: 0.6310
Validation set r-squared scores: 0.0991
Testing set mean squared error: 0.5798
Testing set r-squared scores: 0.1039
```